

Dahua Feng

85 Engineer's Way
Charlottesville, VA, 22903

[Website / GitHub](#)
dahua@email.virginia.edu

Education

B.S. in Information and Computing Sciences, Peking University

Beijing, China

School of Electronics Engineering and Computer Science

Sep. 2020 - Jul. 2024

- **GPA:** 3.700/4.000 (87.4/100)
- **Thesis:** *A Simulator Design for the Storage of Mobile Devices* (supervised by Prof. Jie Zhang)

Ph.D. Student in Computing Science, University of Virginia

Virginia, USA

Department of Computer Science

Aug. 2024 - Present

Publications

1. **Profiling Apple Silicon Performance for ML Training**, Preprint, 2025

Dahua Feng*, Zhiming Xu*, Rongxiang Wang, Felix Xiaozhu Lin

Research Experience

School of Computer Science, Peking University

Mar. 2022 - Jan. 2023

Research Intern

Advisor: Prof. Zhi Yang

Project: DNN acceleration based on graph optimization

- This project focused on the acceleration for neural network computation. Primarily, we tried to use the genetic algorithm based on BFS and DP to schedule the ops. Then we mainly focused on the resources allocation for each of the operations.
- I analyzed part of the source code of Roller and the time evaluation source code of TVM.
- I implemented the functionality of resources allocation for the IOS scheduler and obtained about 10% improvement.

Picasso Lab, University of California, Santa Barbara

Jul. 2023 - Dec. 2023

Research Intern

Advisor: Prof. Yufei Ding

Project: CXL-based memory disaggregated system for DLRM

- This project aimed to improve the performance of the DLRM (deep learning recommendation system) training. We tried to implement a better approach for sharding embedding tables across many GPUs with CXL as memory expansion units.
- I analyzed the problem of the embedding table placement on multi-GPU theoretically and proposed some possible algorithms for the load-balance memory allocation.
- I explored the existing approach to solving sharding problems such as using RL and proposed some potential ways to improve it.

Project: Combination of NVSHMEM with DLRM embedding table lookup

- This project is about to leveraging the unbalanced placement of embedding tables by introducing NVSHMEM as a new inter-GPU communication method. We used CUDA kernel to get realistic data to help us design the strategy.
- I analyzed the realistic data and got some insights about the influence of placement.
- I have completed some important components such as a prediction model to predict the time of table batched embedding.

XSEL Lab, University of Virginia

Aug. 2024-Dec. 2024

Research Assistant

Advisor: Prof. Felix Lin

Project: Apple Silicon for machine learning

- This project focused on the performance of ML training on Apple Silicon. I conducted related experiments and measurements to analyze the performance gap between Apple Silicon and NVIDIA GPUs in ML training, especially for LLM training.
- I conducted the experiments about end-to-end training performance and benchmarked BLAS kernels.
- I organized the result data and made related analysis based on results of measurements.

Teaching Experience

Teaching Assistant of Computer Architectures

Fall 2023

School of Electronic Engineering and Computer Science, Peking University

- As a TA of Computer Architectures course instructed by Prof. Jie Zhang, I participated in the designing and revision of the course projects, conducted Q&A sessions in class, organized the quizzes, and worked on the final exam.

Skills

Programming Languages & Softwares: C, C++, Python, CUDA C

Python Packages: torch, scikit-learn

Languages: Mandarin (native), English (TOEFL iBT: 102/120), Korean (TOPIK 228/300, Level-5)